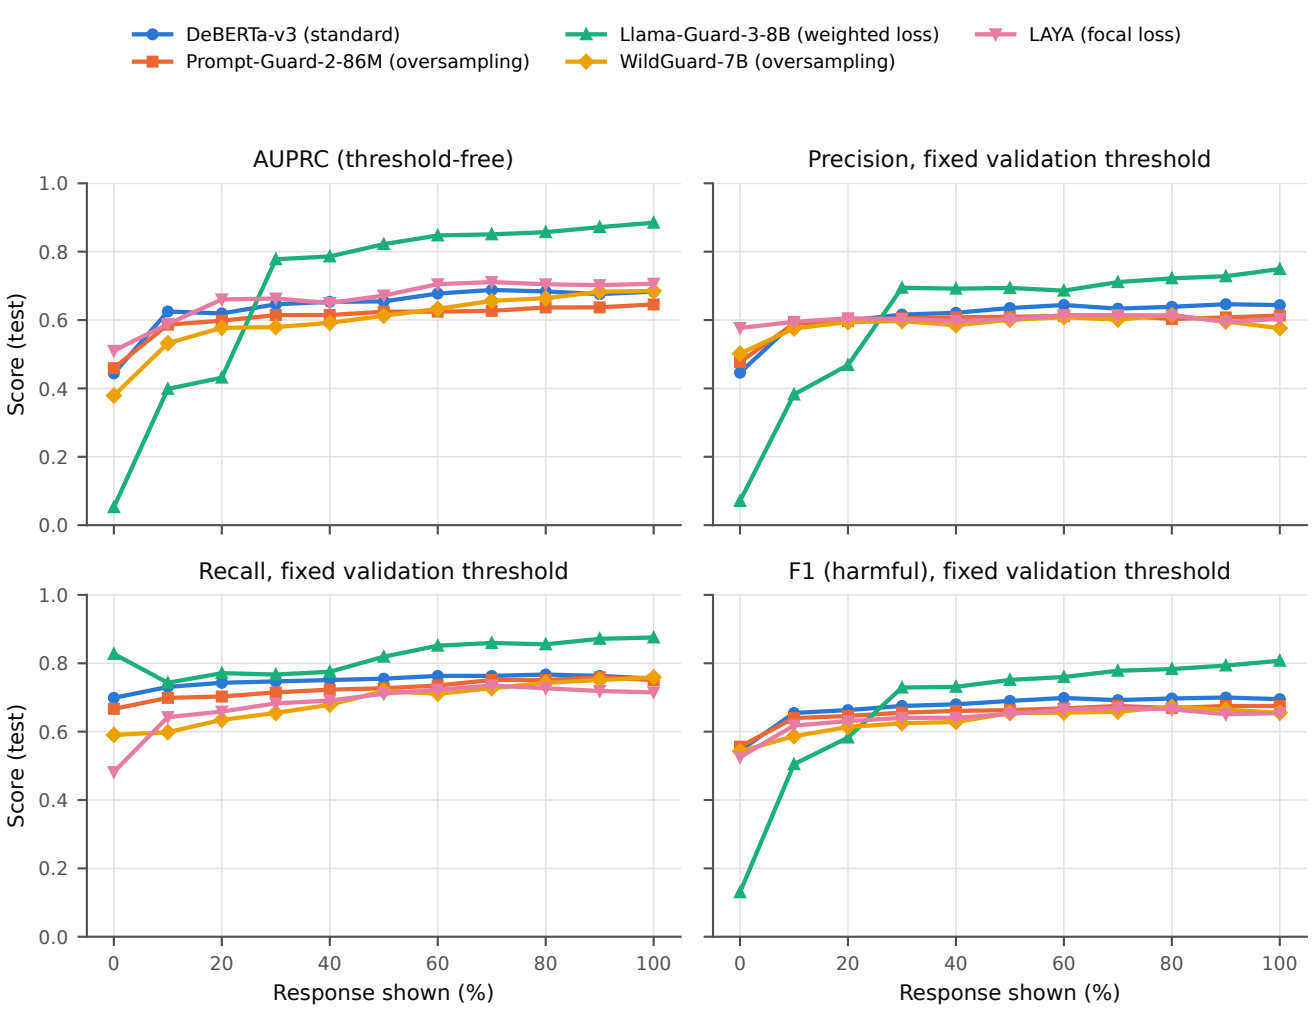
DeBERTa-v3 (standard)
Prompt-Guard-2-86M (oversampling)
Llama-Guard-3-8B (weighted loss)
WildGuard-7B (oversampling)
LAYA (focal loss)
AUPRC (threshold-free)
Precision, fixed validation threshold
Recall, fixed validation threshold
F1 (harmful), fixed validation threshold
Score (test)
Response shown (%)
0.0
0.2
0.4
0.6
0.8
1.0
0
20
40
60
80
100